

LEXINGTON WHALEN

✉ whalenlex[[@](mailto:whalenlex@gmail.com)]gmail.com ·  [lxaw.github.io](https://github.com/lxaw) ·  [linkedin.com/in/lxaw](https://www.linkedin.com/in/lxaw)

PROFILE

Machine learning researcher specializing in efficient training and inference of large language models, with first- and co-first-author publications at ICML, CVPR, and ACL and an NSF Graduate Research Fellowship. Currently co-lead pretraining and inference optimization for SoftBank SB Intuitions' Sarashina LLMs; previously co-led NVIDIA's diffusion language models and led research teams in the U.S. and Japan.

EDUCATION

Georgia Institute of Technology

M.S. in Computer Science, Specialization in Artificial Intelligence

NSF Graduate Research Fellow.

Atlanta, GA

Expected Dec. 2026

University of South Carolina, Honors College

Accelerated B.S./M.S. in Computer Science

Outstanding Senior, Class of 2023.

Columbia, SC

Jan. 2021 – Aug. 2024

PREPRINTS & TECHNICAL REPORTS

- [1] **L. Whalen**, Y. Ito, R. Sakamoto. “Teaching Diffusion to Speculate Left-to-Right.” Technical Report, 2026. [SB Intuitions]
- [2] Y. Fu, **L. Whalen**, A. Garg, C. Wu, M. Khadkevich, N. Oswald, E. Xie, D. Egert, S. Turuvekere Sreenivas, S. Diao, C. Yu, Y. Yu, W. Chen, S. Norouzi, S. Lan, L. Zhu, J. Wang, J. Jiang, M. Mardani, M. Maghoumi, S. Han, A. Jukić, N. Tajbakhsh, J. Kautz, P. Molchanov. “Nemotron-Labs-Diffusion: A Tri-Mode Language Model Unifying Autoregressive, Diffusion, and Self-Speculation Decoding.” Technical Report, 2026. [NVIDIA]

PEER-REVIEWED PUBLICATIONS

- [1] Y. Fu*, **L. Whalen***, Z. Ye, X. Dong, S. Diao, J. Liu, C. Wu, H. Zhang, E. Xie, S. Han, M. Khadkevich, J. Kautz, Y. (Celine) Lin, P. Molchanov. “Efficient-DLM: From Autoregressive to Diffusion Language Models, and Beyond in Speed.” **ICML**, 2026. [NVIDIA]
- [2] C. Shih, C. Li, C. Yu, H. Fang, S. Li, W. Hsin, **L. Whalen**, H. Suh, G. Eisenhauer, L. Liu, Y. (Celine) Lin. “NeRArch-Sim: A Unified Simulator for Benchmarking and DSE of Neural Rendering Accelerators.” **ISCA**, 2026. [Georgia Tech]
- [3] **L. Whalen***, Z. Du*, H. You*, C. Li, S. Li, Y. (Celine) Lin. “Early-Bird Diffusion: Investigating and Leveraging Timestep-Aware Early-Bird Tickets in Diffusion Models for Efficient Training.” **CVPR**, 2025. [Georgia Tech]
- [4] Z. Ye, Z. Wang, K. Xia, J. Hong, L. Li, **L. Whalen**, C. Wan, Y. Fu, Y. (Celine) Lin, S. Kundu. “LAMB: A Training-Free Method to Enhance the Long-Context Understanding of SSMS via Attention-Guided Token Filtering.” **ACL**, 2025. [Georgia Tech]
- [5] A. Murphy, **L. Whalen**, S. Dubinsky, M. Gavin, J. F. Bailyn, J. Ginn. “On ‘historical unity’ of Russian and Ukrainian: A linguistic perspective on language conflict and change.” **LSA**, Apr. 2023. [University of South Carolina]

* Equal contribution

PATENTS

- [1] **L. Whalen**, Y. Ito, R. Sakamoto. 「拡散ベースの投機的ドラフトモデルにおける受理長改善アルゴリズム」 Japanese Patent (to be filed), 2026. [SB Intuitions]
“Algorithm to Improve Acceptance Length of Diffusion-Based Speculative Draft Models.”
- [2] Y. Fu, **L. Whalen**, X. Dong, S. Diao, C. Wu, E. Xie, S. Han, J. Kautz, Y. (Celine) Lin, P. Molchanov. “Training Framework for Converting Autoregressive Language Models into Faster Diffusion Language Models.” U.S. Patent (to be filed), 2026. [NVIDIA]

RESEARCH EXPERIENCE

SoftBank – SB Intuitions

Tokyo, Japan

Pretraining & Inference Optimization Research Lead

Apr. 2026 –

- Co-lead pretraining and post-training optimization research for the Sarashina series, Japan’s largest LLMs, at SoftBank’s SB Intuitions R&D Division.
- Built SB Intuitions’ first [speculative-decoding](#) inference stack for the Sarashina family; introduced training-time drafter-alignment methods that raise the accepted draft length by 21–76% over a position-uniform baseline.
- Serve as a key technical voice in Japanese-language cross-team discussions on pretraining, post-training, and inference optimization, bridging research and engineering.
- Collaborate with leading Japanese and international technology companies and government agencies to evaluate GPU resources and co-develop long-term compute strategies for large-scale training.

NVIDIA

Atlanta, GA & Santa Clara, CA

Efficient Deep Learning Intern & Data Filtering Challenge Lead

Dec. 2024 – Mar. 2026

Mentors: Dr. Pavlo Molchanov & Yonggan Fu (NVIDIA Research).

- Co-developed [Nemotron-Labs-Diffusion](#), a tri-mode (autoregressive / diffusion / self-speculation) language model family at 3B–14B scale; the 8B model decodes $6\times$ more tokens per forward pass than Qwen3-8B at comparable accuracy (up to $4\times$ higher throughput); [released on Hugging Face](#) with $\sim 1\text{M}$ total downloads.
- Co-first author of [Efficient-DLM](#) (ICML 2026), converting pretrained autoregressive models into diffusion LMs; the 8B model reaches $+5.4\%/+2.7\%$ higher accuracy at $4.5\times/2.7\times$ higher throughput than Dream 7B / Qwen3 4B (open-weights on Hugging Face); presented directly to CEO Jensen Huang.
- Built the training, evaluation, and inference infrastructure adopted by 5+ internal research teams for vision-language models, embeddings, and latent reasoning.
- Organized NVIDIA’s international data filtering challenge across 30+ university and industry teams, coordinating GPU resource allocation with sponsors including Lambda Labs and Turing.

Georgia Institute of Technology

Atlanta, GA

Graduate Researcher

Aug. 2024 – Dec. 2025

- Led a team of 3 Ph.D. and Masters students (co-first author) to cut diffusion-model training time by $2.9\text{--}5.8\times$ (up to $10.3\times$ vs. standard train-prune-finetune) while maintaining generation quality; [CVPR 2025](#).
- Algorithm lead for the United States’ ARPA-H THOR healthcare initiative: reduced cancer detection sensor degradation by over 70%, extending sensor lifetime from a few hours to several days.
- Engineered resource-efficient ML architectures powered by energy from the human body, collaborating with teams from Stanford, Northwestern, Carnegie Mellon, and MIT.
- Presented technical progress to federal sponsors and healthcare stakeholders nationwide.

HONORS & AWARDS

- | | |
|--|-----------|
| • SoftBank Group Representative, 2026 Entrance Ceremony | Mar. 2026 |
| • NSF Graduate Research Fellowship (\$150,000 value) | Apr. 2024 |
| • Toyo University Exchange Student Representative | Jul. 2023 |
| • University of South Carolina Outstanding Senior, Class of 2023 | Feb. 2023 |
| • U.S. Department of State Critical Language Scholar of Japanese | Jan. 2023 |
| • University of South Carolina Top Scholar (\$100,000+ value) | Aug. 2020 |

INVITED TALKS

- | | |
|---|-----------|
| • “拡散言語モデル：基本から応用まで”, SoftBank – SB Intuitions | Jul. 2026 |
| <i>Diffusion Language Models: From Fundamentals to Applications</i> | |
| • 入社式スピーチ, SoftBank | Apr. 2026 |
| <i>Entrance Ceremony Speech</i> | |

LANGUAGE SKILLS

English (Native) · Japanese (Native Level, JLPT N1 Certified) · Mandarin (Intermediate)